

Topic 11: Lasso And Beyond: Convex Learning

by Sai Zhang

Key points:

Disclaimer:

11.1 Lasso

Lasso (Least absolute Shrinkage and Selection Operator), proposed by Tibshirani (1996), aims to minimize the SSR (sum of residual squares) subject to the L1-norm (sum of the absolute value) of the coefficients being less than a constant.

11.1.1 Set up

For data $(\mathbf{x}_i, y_i)_{i=1}^n$, where

- y_i is the outcome for individual i
- $\mathbf{x}_i = (x_{i1}, \dots, x_{ip})'$ is the $p \times 1$ vector of predictors

Then the Lasso estimator $(\hat{\alpha}, \hat{\beta})$ is defined as

$$(\hat{\alpha}, \hat{\beta}) = \arg \min_{\alpha, \beta} \left\{ \sum_{i=1}^n \left(y_i - \alpha - \sum_{j=1}^p \beta_j x_{ij} \right)^2 \right\} \quad \text{s.t.} \quad \sum_{j=1}^p |\beta_j| \leq t$$

for the $n \times 1$ response vector $\mathbf{y} = (y_1, \dots, y_n)'$, the $n \times p$ design matrix $\mathbf{X} = (\mathbf{x}_1, \dots, \mathbf{x}_n)'$ where $\mathbf{x}_i = (x_{i1}, \dots, x_{ip})'$ is a $p \times 1$ vector. Here $\hat{\alpha} = \bar{y}$, w.l.o.g., let $\bar{y} = 0$ and omit α for simplicity.

In matrix form, we have

- constrained form:

$$\hat{\beta} = \arg \min_{\beta \in \mathbb{R}^p} \left\{ \frac{1}{2n} \|\mathbf{y} - \mathbf{X}\beta\|_2^2 \right\} \quad \text{s.t.} \quad \|\beta\|_1 \leq t$$

- unconstrained form:

$$\hat{\beta}(\lambda) = \arg \min_{\beta \in \mathbb{R}^p} \left\{ \frac{1}{2n} \|\mathbf{y} - \mathbf{X}\beta\|_2^2 + \lambda \|\beta\|_1 \right\}$$

where the regularization parameter $\lambda \geq 0$:

- $\lambda \rightarrow \infty$: $\hat{\beta}_{lasso} = \mathbf{0}$
- $\lambda = 0$: $\hat{\beta}_{lasso} \rightarrow \hat{\beta}_{OLS}$


11.1.2 Solving Lasso

Lasso is essentially a quadratic optimization problem. Hence, the solution is given by taking the derivative (of the unconstrained question) and set it equal to 0

$$\frac{d}{d\beta} \left(\frac{1}{2n} \|\mathbf{y} - \mathbf{X}\beta\|_2^2 + \lambda \|\beta\|_1 \right) = 0$$

$$\Rightarrow \underbrace{\frac{1}{n} \mathbf{X}'}_{p \times n} \underbrace{(\mathbf{y} - \mathbf{X}\beta)}_{= \epsilon, n \times 1} = \lambda \begin{cases} \text{sign}(\beta_j), & \beta_j \neq 0 \\ [-1, 1], & \beta_j = 0 \end{cases}$$

this result follows the fact the L-1 norm $\|\beta\|_1$ is piecewise linear:



L1-norm (1-dimension)

For each component of the vector of the L-1 norm $f(\beta_j) = |\beta_j|$, we have:

- $\beta_j > 0$: $f'(\beta_j) = 1$
 - $\beta_j < 0$: $f'(\beta_j) = -1$
 - $\beta_j = 0$: $df \in [-1, 1]$ (shaded area)
- which gives the results stated above.

Take another look at this result

Proposition 11.1.1: Lasso Parameter Selection Rule

$$\frac{1}{n} \mathbf{X}' (\mathbf{y} - \mathbf{X}\beta) = \frac{1}{n} \mathbf{X}' \epsilon = \lambda \begin{cases} \text{sign}(\beta_j), & \beta_j \neq 0 \\ [-1, 1], & \beta_j = 0 \end{cases}$$

which gives a parameter selection criterion: for $\beta_j \neq 0$, $\text{sign}(\beta_j)$ **must agree** with $\text{Corr}(\mathbf{x}_j, \epsilon)$, the correlation between the j -th variable \mathbf{x}_j and (full-model) residuals $\epsilon = \mathbf{y} - \mathbf{X}\beta$.

11.1.3 Algorithm: LARS


Mathematically, Lasso is quite intuitive, but computationally, it can be quite consuming. Efron et al. (2004) propose an algorithm that takes steps from a all-0 model to the biggest model (OLS), that is, **Least Angle Regression (LARS)**.

Intuition

The basic intuition of LARS is quite straight-forward: covariates are considered from the **highest** correlation with \mathbf{y} (*smallest* angle from \mathbf{y}) to the **least** correlated one (*largest* angle from \mathbf{y}) (illustrated below).

And the steps of the LARS algorithm are

- 1 start with the null model $\hat{\beta} = \mathbf{0}$: $\hat{\mu} = \mathbf{X}'\mathbf{0} = \mathbf{0}$



- 2 calculate residual vector $\mathbf{r} = \mathbf{y} - \hat{\boldsymbol{\mu}}$
- 3 determine the correlation vector between \mathbf{r} and each parameter $\mathbf{x}_j, \forall j = 1, \dots, p: \mathbf{X}'\mathbf{r}$
- 4 pick the largest correlation $\mathbf{x}_{\text{step1},1}^*$, increase its $\hat{\beta}$ to the point where its correlation with \mathbf{r} will be **equal** with that of another parameter $\mathbf{x}_{\text{step1},2}^*$
- 5 next, increase the $\hat{\beta}$ for both $\mathbf{x}_{\text{step1},1}^*, \mathbf{x}_{\text{step1},2}^*$ in an **equiangular** direction between these two, until a third parameter becomes equally important

And keep looping this way, until all the predictors enter the model and eventually $\mathbf{X}'\mathbf{r} = 0$

Properties of LARS

LARS has several properties:

- geometrically travels in the direction of **equal** angle to all active covariates
- assume all covariates are independent
- computationally quick: only take m steps, where m is the number of parameters being considered

And it is in between 2 classic model-selection methods: **Forward Selection** and **Stagewise Selection**:

- **Forward Selection**

- for \mathbf{y} , select the most correlated \mathbf{x}_{j_1}
- regress \mathbf{x}_{j_1} on \mathbf{y} , get the residuals
- select the most correlated \mathbf{x}_{j_2} with the residual of \mathbf{y} net of \mathbf{x}_{j_1}

looping this, for a k -parameter linear model, it takes k steps. Forward Selection is an aggressive fitting technique, can be overly greedy (some important predictors may be eliminated due to correlation with already selected variables).

- **Forward Stagewise**

- also begin with $\hat{\boldsymbol{\mu}} = 0$
- for a current Stagewise estimate $\hat{\boldsymbol{\mu}}$, the current residual vector is then $\mathbf{y} - \hat{\boldsymbol{\mu}}$, its correlation with \mathbf{X} is then $\mathbf{X}'(\mathbf{y} - \hat{\boldsymbol{\mu}}) \equiv \hat{\mathbf{c}}$

- next, heavily computational, go in the direction of the greatest current correlation, but by only a **small** step

$$\hat{j} = \arg \max |\hat{c}_j|, \hat{\mu} \rightarrow \hat{\mu} + \epsilon \cdot \text{sign}(\hat{c}_{\hat{j}}) \cdot \mathbf{x}_{\hat{j}}$$

here, ϵ is a **small** constant, hence avoiding the greediness of Forward Selection, at a cost of computational efficiency¹.

LARS avoids the over-greediness of Forward Selection and computational heaviness of Forward Stagewise.

11.1.4 From LARS to Lasso

The Lasso algorithm is built upon LARS, with the constraint from the mathematical condition of Proposition

11.1.1: $\text{sign}(\beta_j)$ must agree with $\text{Corr}(\mathbf{x}_j, \epsilon)$.

Theorem 11.1.2: Lasso Modification Condition

If $\tilde{\gamma} < \hat{\gamma}$, stop the ongoing LARS step at $\gamma = \tilde{\gamma}$ and remove j from the calculation of the next equiangular direction, where

- the path at any LARS step is

$$\beta(\gamma), \beta_j(\gamma) = \hat{\beta}_j + \gamma \hat{d}_j$$

\hat{d}_j specifies the **direction** to take the j -th component, γ is **how far** to travel in the direction of \hat{d}_j before adding in a new covariate

- $\hat{\gamma}$ represents the smallest **positive** value of γ s.t. some new covariate joins the active set (the set of covariates used on path)
- $\tilde{\gamma}$ represents the first time $\beta_j(\gamma)$ **changes signs**.

The key point of 11.1.2 is that Lasso does **NOT** allow the $\hat{\beta}_j$ to change signs, if it changes sign, it will be subtracted from the active set. Now, from this point of view, we can compare the 3 algorithms:

LARS	no sign restrictions
Lasso	$\hat{\beta}_j$ agrees in sign with \hat{c}_j
Stagewise	successive differences of $\hat{\beta}_j$ agree in sign with the current correlation $\hat{c} = \mathbf{x}'_j(\mathbf{y} - \hat{\mu})$

Again, LARS requires the least steps but is most greedy, Stagewise is computationally consuming but robust. Lasso is in between.

11.2 Penalized Least Square Estimation

Lasso is one special class of Penalized Least Square (PLS) Estimation. For the linear regression model $\mathbf{y} = \mathbf{X}\beta + \epsilon$, if $\epsilon \sim \mathcal{N}(0, \sigma^2 \mathbf{I}_n)$, we have PLS as

$$\min_{\beta \in \mathbb{R}^p} \left\{ \frac{1}{2n} \|\mathbf{y} - \mathbf{X}\beta\|_2^2 + \sum_{j=1}^p p_\lambda(|\beta_j|) \right\}$$



¹Forward Selection is essentially choosing $\epsilon = |\hat{c}_{\hat{j}}|$

where $p_\lambda(\cdot)$ is a penalty function indexed by the regularization parameter $\lambda \geq 0$. **Antoniadis and Fan (2001)** showed that the PLS estimator $\hat{\beta}$ has the following properties:

- **sparsity**: if $\min_{t \geq 0} \{t + p'_\lambda(t)\} > 0$
- **approximate unbiasedness**: if $p'_\lambda(t) = 0$ for t large enough
- **continuity**: iff $\arg \min_{t \geq 0} \{t + p'_\lambda(t)\} = 0$

In general

- the **sigularity** of penalty function at the origin, $p'_\lambda(0_+) > 0$ is needed for generating **sparsity** in variable selection
- the **concavity** is needed to reduce the bias



References

- Anestis Antoniadis and Jianqing Fan. Regularization of wavelet approximations. *Journal of the American Statistical Association*, 96(455):939–967, 2001.
- Bradley Efron, Trevor Hastie, Iain Johnstone, and Robert Tibshirani. Least angle regression. *The Annals of Statistics*, 32(2):407 – 499, 2004. doi: 10.1214/0090536040000000067. URL <https://doi.org/10.1214/0090536040000000067>.
- Robert Tibshirani. Regression shrinkage and selection via the lasso. *Journal of the Royal Statistical Society: Series B (Methodological)*, 58(1):267–288, 1996.